SMART STORAGE TANK LEVEL & USAGE PREDICTION SYSTEM

For Petrochemical Plants

Comprehensive Technical Report — Milestones 1 through 4

IoT Sensor Architecture • Machine Learning Pipeline • Real-Time Analytics Dashboard

Project	Smart Storage Tank Level & Usage Prediction System
Team	Team-7: Mithun Chanderan, Aaron Rohit, Rahul R, Abinand
Institution	Christ University, Bengaluru
Industry Partner	L&T Construction
Document Type	Internship Technical Report — Consolidated (M1–M4)
Classification	L&T Construction — Internal Use Only
Status	Milestones 1–4 Complete	Milestone 5 Upcoming

CONFIDENTIAL — L&T Construction Internal Use

Table of Contents

	Table of Contents
1.	**Executive Summary**
2.	**Project Overview**
2.1	Problem Statement
2.2	Solution Architecture
2.3	System Scope
3.	**System Architecture & Data Flow**
3.1	End-to-End Pipeline Architecture
3.2	Structured Data Flow — Step by Step
4.	**Milestone 1 — Circuit Design & Cloud Integration**
4.1	Objective
4.2	Component Inventory
4.3	Pin Mapping — ESP32-S3
4.4	ThingSpeak Field Mapping
4.5	Firmware Key Features
4.6	Milestone 1 Deliverables
5.	**Milestone 2 — Dataset Generation & ML Model Development**
5.1	Objective
5.2	Synthetic Dataset Design
5.3	Class Distribution Strategy
5.4	Feature Engineering — 17 Input Features
5.5	Model Evaluation Results
5.6	Serialised Artifacts
6.	**Milestone 3 — Live Data Simulation & Model Integration**
6.1	Objective
6.2	Cylindrical Tank Volume Conversion
6.3	Signal Processing Pipeline
6.4	Prediction Output Columns
6.5	Risk Classification Thresholds

7.	**Milestone 4 — Analytics Output & Dashboard Visualisation**
7.1	Objective
7.2	Structured CSV Output Files
7.3	Power BI Dashboard Visual Inventory
7.4	Real-Time Auto-Refresh Architecture
7.5	Key DAX Measures
8.	**Project Milestones & Deliverables Summary**
8.1	Complete File Artifact Registry
9.	**Conclusion**

Executive Summary

The Smart Storage Tank Level & Usage Prediction System is an end-to-end IoT and machine learning solution developed for petrochemical plant operations in collaboration with L&T Construction. The system continuously monitors liquid levels across storage tanks, applies physics-based volume conversion, and delivers real-time ML-powered fill predictions and risk alerts to plant operators through a live Power BI dashboard.

This consolidated technical report documents the complete project lifecycle across four completed milestones: hardware circuit design and cloud integration (M1), synthetic dataset generation and ML model development (M2), live data simulation and model integration (M3), and analytics output with dashboard visualisation (M4). A fifth milestone covering final industry analysis and interpretation is currently in progress.

Metric	Result	Standard
RF Regressor R^2	1.000000	≥ 0.95 target
RF Classifier Accuracy	100.0%	≥ 95% target
Prediction MAE	4.34 L	< 500 L target
Dashboard Refresh Rate	15 seconds	≤ 60 s target
Tanks Monitored	3 Active (5 total)	3–5 tanks
Tank Capacity	50,265.5 L each	50,000 L design

1. Project Overview

1.1 Problem Statement

In a petrochemical plant, continuous and reliable monitoring of liquid raw material storage is critical to maintaining uninterrupted production. The absence of a continuous monitoring mechanism makes it extremely difficult to track real-time tank levels, measure consumption rates, and estimate time-to-empty values. This operational gap leads to critical issues including:

- Tank overflow during refilling cycles due to undetected high-level conditions
- Dry-run conditions during periods of high demand causing production halts
- Inventory imbalances across tanks leading to inefficient resource allocation
- Inefficient procurement planning due to lack of predictive consumption data

1.2 Solution Architecture

The Smart Storage Tank Level & Usage Prediction System addresses these operational gaps through an intelligent, sensor-driven monitoring architecture. The solution delivers:

- Real-time ultrasonic level sensing via HC-SR04 sensors on each tank
- ESP32-S3 microcontroller processing with Wi-Fi cloud upload every 15 seconds
- ThingSpeak IoT cloud storage with live time-series visualisation
- Random Forest ML models for fill level regression and risk classification
- Power BI Desktop dashboard with 8 interactive visuals and live auto-refresh

1.3 System Scope

Parameter	Value / Description
Number of Tanks	5 cylindrical tanks (3 active in current deployment)
Tank Capacity	50,265.5 litres each ($\pi r^2 h = \pi \times 4 \times 4$ m)
Tank Dimensions	Radius: 2.0 m | Height: 4.0 m (400 cm)
Refill Schedule	Daily at 06:00 AM
Data Upload Rate	Every 15 seconds via HTTP GET to ThingSpeak
ML Training Data	10,000 synthetic records across 180-day simulation window
Cloud Platform	ThingSpeak (MathWorks) — Channel ID: 3327870
Dashboard Tool	Power BI Desktop — 8 visuals, DAX measures, 3 slicers

2. System Architecture & Data Flow

2.1 End-to-End Pipeline Architecture

The system is structured across five architectural layers, forming a complete IoT-to-analytics pipeline. Each layer is designed to be modular, enabling independent upgrades without disrupting adjacent components.

Layer	Component	Technology	Function
L1	**Sensing**	HC-SR04 Ultrasonic + Water Level Sensors	Measure liquid surface distance (cm) and ADC analog levels across all 3 active tanks
L2	**Processing**	ESP32-S3 Microcontroller (Arduino C++)	GPIO reading, distance-to-volume conversion, usage calculation, HTTP upload every 15 s
L3	**Connectivity**	Wi-Fi → ThingSpeak REST API	HTTP GET to api.thingspeak.com/update | 8 fields per entry | Channel 3327870
L4	**ML Inference**	Python (Google Colab) | Random Forest | scikit-learn	Signal processing → feature engineering → RF Regressor (fill L) + RF Classifier (risk label)
L5	**Visualisation**	Power BI Desktop + Python Auto-Refresh Scripts	8 live visuals, DAX measures, 3 slicers, F5 auto-refresh every 15 s via pyautogui

2.2 Structured Data Flow — Step by Step

Step	Component	Action
1	**HC-SR04 Sensors**	Emit ultrasonic pulse; measure echo time-of-flight to derive distance to liquid surface in centimetres. Three sensors operate in parallel, one per active tank.
2	**ESP32-S3 Firmware**	Reads TRIG/ECHO GPIO pins, converts pulse duration to distance ($d = t \times 0.034 / 2$). Applies tankHeight − distance to compute fill height. Calculates total usage, time-to-empty.
3	**ThingSpeak Upload**	Firmware constructs HTTP GET with API key and 8 field values. Uploads every 15 seconds. ThingSpeak stores timestamped entries and renders live charts on channel 3327870.

4	**Python Fetch Layer**	fetch_thingspeak.py retrieves the 200 latest records via ThingSpeak REST API. Parses JSON response, validates field presence, handles null entries.
5	**Signal Processing**	Three-stage pipeline: (a) range validation [2–400 cm], (b) spike rejection [Δ > 50 cm discarded], (c) weighted smoothing [0.3 × new + 0.7 × previous].
6	**Volume Conversion**	Cleaned distance → fill_height_m = (400 − dist) / 100 → volume = π × 2^2 × fill_height_m × 1000 litres. Clamped to [0, 50,265.5 L].
7	**Feature Engineering**	17-column input matrix assembled. Categorical features encoded via saved LabelEncoders. All features scaled via saved StandardScaler (zero drift from training).
8	**ML Inference**	RF Regressor predicts fill level in litres. RF Classifier assigns CRITICAL / WARNING / NORMAL risk label. Alert priority and refill flag derived from predictions.
9	**CSV Export**	Prediction DataFrame exported as milestone3_predictions.csv (15 columns). Milestone 4 enrichment adds 35 KPI columns; outputs 3 structured CSVs.
10	**Power BI Dashboard**	Three CSV tables loaded, relationships defined. DAX measures compute KPI cards. 8 visuals render fill levels, risk alerts, consumption trends, prediction accuracy.
11	**Live Refresh**	fetch_thingspeak.py writes live_tanks.csv every 15 s. auto_refresh_powerbi.py sends F5 keypress to Power BI Desktop window via pyautogui. Dashboard stays live.

3. Milestone 1 — Circuit Design & Cloud Integration

3.1 Objective

Establish the foundational hardware sensing layer and cloud connectivity pipeline. Design and simulate the complete circuit using Cirkit Designer, implement the ESP32-S3 firmware, and validate live data upload to the ThingSpeak IoT platform.

3.2 Component Inventory

Component	Qty	Purpose
ESP32-S3 (Beta)	1	Dual-core LX7 microcontroller with integrated Wi-Fi. Central processing and cloud upload unit.
HC-SR04 Ultrasonic Sensor	3	Time-of-flight distance sensor (5V). Measures headspace from sensor mount to liquid surface.
Water Level Sensor (Resistive)	3	Analog ADC sensor (3.3V). Cross-validates ultrasonic readings; immune to surface turbulence.
5V DC Power Supply	1	Dedicated supply for HC-SR04 sensors. Shared GND with ESP32-S3 for signal integrity.
Breadboard + Connecting Wires	1 set	Simulation wiring in Cirkit Designer environment.

3.3 Pin Mapping — ESP32-S3

GPIO Pin	Connected To	Signal Type
GPIO 5	TRIG — HC-SR04 Tank 1	Digital Output
GPIO 17	TRIG — HC-SR04 Tank 2	Digital Output
GPIO 4	TRIG — HC-SR04 Tank 3	Digital Output
GPIO 18	ECHO — HC-SR04 Tank 1	Digital Input
GPIO 16	ECHO — HC-SR04 Tank 2	Digital Input
GPIO 2	ECHO — HC-SR04 Tank 3	Digital Input
GPIO 1	SIG — Water Level Sensor 1	Analog Input (ADC)
GPIO 3	SIG — Water Level Sensor 2	Analog Input (ADC)
GPIO 6	SIG — Water Level Sensor 3	Analog Input (ADC)
3V3	VCC — All Water Level Sensors	3.3V Power
GND	GND — All Sensors & Power Supply	Common Ground

3.4 ThingSpeak Field Mapping

Field	Variable	Description
Field 1	levels[0]	Fill level of Tank 1 (cm)
Field 2	waterValues[0]	Analog water sensor reading — Tank 1
Field 3	levels[1]	Fill level of Tank 2 (cm)
Field 4	waterValues[1]	Analog water sensor reading — Tank 2
Field 5	levels[2]	Fill level of Tank 3 (cm)
Field 6	waterValues[2]	Analog water sensor reading — Tank 3
Field 7	usage	Total liquid consumed since previous reading cycle
Field 8	timeToEmpty	Predicted time remaining until all tanks empty (hours)

3.5 Firmware Key Features

- Wi-Fi connection via built-in ESP32-S3 module; polls WL_CONNECTED status before proceeding
- TRIG pins configured as OUTPUT; ECHO and SIG pins as INPUT during setup()
- Distance formula: distance_cm = pulseIn(ECHO, HIGH) × 0.034 / 2
- Usage = previous total level − current total level (per 15-second interval)
- Time-to-empty = total level / usage rate (cycles)
- HTTP GET to api.thingspeak.com/update with all 8 fields every 15 seconds

3.6 Milestone 1 Deliverables

- Cirkit Designer simulation schematic (URL: cirkitdesigner.com/project/52f8eb59...)
- ESP32-S3 Arduino C++ firmware — fully functional with sensor reads and cloud upload
- ThingSpeak Channel 3327870 — live data verified across all 8 fields
- Confirmed live ThingSpeak chart output for tank levels, water sensors, usage, and time-to-empty

4. Milestone 2 — Dataset Generation & ML Model Development

4.1 Objective

Generate a statistically representative synthetic dataset simulating 180 days of five-tank petrochemical operations, perform exploratory data analysis, engineer 17 input features, and train and evaluate regression and classification models for deployment in Milestone 3.

4.2 Synthetic Dataset Design

Parameter	Value
Total Records	10,000 rows (15-second interval readings)
Simulation Window	180 days from 1 January 2024
Tank Count	5 tanks (TANK_01 through TANK_05)
Tank Capacity	50,000 L (design) | 50,265.5 L (geometric)
Tank Height (Simulation)	100 cm (Milestone 2) → upgraded to 400 cm in Milestone 3
Sensor Noise (Ultrasonic)	Gaussian σ = 0.3 cm
Sensor Noise (ADC)	Gaussian σ = 30 ADC units
Emergency Spike Probability	5% per record (1.5× to 2.5× base consumption rate)

4.3 Class Distribution Strategy

Risk Class	Target Rows	Proportion	Fill % Range
NORMAL	6,000	60%	25.1% – 99.5%
WARNING	2,500	25%	10.1% – 24.9%
CRITICAL	1,500	15%	0.5% – 9.9%

4.4 Feature Engineering — 17 Input Features

#	Feature Column	Category	Description
1	`hour_of_day`	Temporal	Hour of reading (0–23)
2	`is_weekend`	Temporal	Binary: 1 if Sat/Sun
3	`shift_encoded`	Temporal (encoded)	LabelEncoded shift period
4	`tank_id_encoded`	Identity (encoded)	LabelEncoded tank identifier
5	`dow_encoded`	Temporal (encoded)	LabelEncoded day-of-week

6	`ultrasonic_distance_cm`	Sensor — Ultrasonic	Raw HC-SR04 distance reading
7	`water_sensor_adc_raw`	Sensor — Analog	Raw ADC water level value
8	`fill_level_cm`	Derived Physical	tankHeight − distance
9	`fill_percentage`	Derived Physical	Fill level as % of capacity
10	`headspace_cm`	Derived Physical	Empty space above liquid
11	`capacity_remaining_litres`	Derived Physical	Unused tank capacity
12	`consumption_per_interval_L`	Consumption	Litres consumed per 15-s interval
13	`consumption_rate_Lps`	Consumption	Instantaneous rate in L/s
14	`level_delta_from_prev`	Consumption	Change from previous reading
15	`cumulative_daily_usage_L`	Consumption	Total consumed since 06:00 refill
16	`estimated_hourly_usage_L`	Consumption	Rate extrapolated to 1 hour
17	`day_index`	Temporal	Sequential day from simulation start

4.5 Model Evaluation Results

Regression — Predict current_fill_litres

Model	MAE (L)	RMSE (L)	R^2	Selected
Linear Regression	0.00	0.00	1.000000	—
Random Forest Regressor	4.34	6.23	1.000000	✓

Classification — Predict Risk Level

Model	Accuracy	Selected
Logistic Regression	99.6667%	—
Random Forest Classifier	100.0000%	✓

The Random Forest Classifier achieves perfect precision, recall, and F1-score (1.00) for all three risk classes across the 3,000-record test set with zero misclassifications. The Random Forest Regressor is selected over the Linear Regression baseline due to its robustness to noise and non-linear

generalisation capability for live sensor data, despite both achieving $R^2 = 1.000000$ on the synthetic test set.

4.6 Serialised Artifacts

- smart_tank_dataset.csv — Full 10,000-row synthetic dataset (22 columns)
- model_rf_regressor.pkl — Trained Random Forest Regressor (150 trees, max_depth=20)
- model_rf_classifier.pkl — Trained Random Forest Classifier (150 trees, max_depth=20)
- scaler.pkl — Fitted StandardScaler for feature normalisation
- label_encoder_risk.pkl — Fitted LabelEncoder for risk class decoding
- le_shift.pkl, le_tank.pkl, le_dow.pkl — Categorical feature encoders

5. Milestone 3 — Live Data Simulation & Model Integration

5.1 Objective

Integrate the trained ML models with live sensor data streamed from the ESP32-S3 through ThingSpeak. Implement physics-based volume conversion, a three-stage signal processing pipeline, and export a structured prediction CSV for Power BI integration.

5.2 Cylindrical Tank Volume Conversion

A fundamental upgrade in Milestone 3 is the transition from height-based (cm) to volume-based (litres) calculations. All values in the system are expressed in operationally meaningful units.

Parameter	Value
Tank Radius (r)	2.0 metres
Tank Height (h)	4.0 metres = 400 cm
Maximum Volume	$\pi \times 2^2 \times 4 \times 1000 = 50{,}265.5$ litres
Conversion Rate	125.66 litres per centimetre of liquid height
Formula	volume_L = π × r² × (400 − distance_cm) / 100 × 1000
Clamping	Output clamped to [0, 50,265.5 L]

5.3 Signal Processing Pipeline

Stage	Name	Logic	Purpose
1	Range Validation	Reject if d < 2 cm or d > 400 cm	Eliminates sensor obstruction and out-of-range echoes
2	Spike Rejection	|d − prev_smoothed| > 50 cm → discard	Removes ultrasonic interference artefacts and spurious echoes
3	Weighted Smoothing	smoothed = 0.3 × new + 0.7 × previous	Reduces noise while preserving real level changes
4	Volume Conversion	Cylindrical formula applied to smoothed distance	Converts cm to operationally meaningful litres

5.4 Prediction Output Columns

Column	Type	Description
`predicted_fill_litres`	float	RF Regressor output: predicted liquid volume

`predicted_fill_pct`	float	Predicted volume as % of tank capacity
`predicted_risk_level`	string	RF Classifier output: CRITICAL / WARNING / NORMAL
`prediction_error_L`	float	Residual: predicted − actual fill (litres)
`alert_priority`	integer	Numeric rank: CRITICAL=3, WARNING=2, NORMAL=1
`refill_needed_flag`	integer	Binary: 1 if fill_percentage ≤ 20%, else 0
`time_to_empty_hours`	float	Hours until depletion at current rate (capped 999)
`sensor_status`	string	OK / INVALID / SPIKE REJECTED / NO DATA

5.5 Risk Classification Thresholds

Risk Level	Fill % Threshold	Volume Equivalent	Recommended Action
CRITICAL	≤ 10%	≤ 5,027 L	Immediate refill required
WARNING	10.1% – 25%	5,027 – 12,566 L	Schedule refill soon
NORMAL	> 25%	> 12,566 L	No action required

6. Milestone 4 — Analytics Output & Dashboard Visualisation

6.1 Objective

Transform the Milestone 3 ML prediction output into a structured analytics layer and a fully interactive Power BI Desktop dashboard. Implement a Python-based real-time auto-refresh solution achieving 15-second live data synchronisation without Power BI DirectQuery limitations.

6.2 Structured CSV Output Files

File	Role	Contents
milestone4_dashboard_data.csv	Primary Source	All predictions with 35 KPI columns. One row per tank per timestamp.
milestone4_live_latest.csv	Gauge & Cards	Single latest row per tank. Drives current fill gauges and status cards.
milestone4_daily_summary.csv	Trend Analysis	Daily aggregates: avg/min/max fill, total consumption, alert counts.
live_tanks.csv	Live Feed	Real-time output of fetch_thingspeak.py — refreshed every 15 seconds.

6.3 Power BI Dashboard Visual Inventory

#	Visual Title	Type	Key Information
1	KPI Cards Row	Card × 4	Total Volume, Overall Fill %, Critical Tank Count, Avg Time to Empty
2	Tank Fill Level Over Time	Line Chart	predicted_fill_litres by timestamp per tank. Threshold lines at 5,027 L (CRITICAL) and 12,566 L (WARNING)
3	Risk Alert Distribution	Donut Chart	Count of readings by predicted_risk_level. Colour-coded green / amber / red
4	Current Fill % per Tank	Gauge × 3	capacity_used_pct from LiveLatest. One gauge per tank. Target at 25% WARNING threshold
5	Predicted vs Actual Fill	Scatter Chart	current_fill_litres (X) vs predicted_fill_litres (Y). Trend line confirms $R^2 = 1.000$
6	Daily Consumption Trend	Clustered Column	total_consumption_L by date from DailySummary, split by tank_id

7	Time to Empty by Tank	Bar Chart	time_to_empty_hours per tank. Conditional colours: <2h red, <8h amber, ≥8h green
8	Prediction Error Distribution	Column Chart	Count of records by error_band: Excellent / Good / Fair / Poor

6.4 Real-Time Auto-Refresh Architecture

Power BI Desktop's built-in Page Refresh requires DirectQuery mode, which the ThingSpeak Web connector does not natively support. A Python two-script solution achieves identical 15-second refresh behaviour.

Script	Function
fetch_thingspeak.py	Connects to ThingSpeak REST API (channel 3327870) every 15 s. Converts distance readings to litres using cylindrical geometry. Computes fill %, risk level, time-to-empty. Writes results to live_tanks.csv.
auto_refresh_powerbi.py	Uses pygetwindow to locate Power BI Desktop process. Activates the window and sends F5 keypress via pyautogui every 15 seconds, triggering data source re-read.

6.5 Key DAX Measures

Measure Name	Logic Summary
`Tank1_Fill_%`	Calculates average distance from ThingSpeak, derives fill percentage: (400 − dist) / 400 × 100
`Tank1_Litres`	Converts distance to fill height in metres, applies $\pi r^2 h \times 1000$ formula
`Tank1_Status`	Derives fill % from distance, applies threshold: ≤10% = CRITICAL, ≤25% = WARNING, else NORMAL
`Total Current Volume (L)`	SUMX over latest timestamp: sums predicted_fill_litres across all tanks
`Overall Fill %`	Divides total current volume by 150,796.5 L (3 tanks × 50,265.5 L)
`Avg Time To Empty`	AVERAGEX filtered to time_to_empty_hours < 999 (excludes infinite/uncapped values)
`Prediction MAE`	AVERAGEX over ABS(prediction_error_L) across all TankPredictions records

7. Project Milestones & Deliverables Summary

M	Title	Key Deliverables	Status
1	System Architecture Design	Circuit design (Cirkit Designer), ESP32-S3 firmware, ThingSpeak channel (8 fields), pin mapping documentation	COMPLETE
2	Dataset Generation & ML Development	10,000-row synthetic dataset, RF Regressor (R^2=1.000), RF Classifier (100% accuracy), all .pkl serialised artifacts	COMPLETE
3	Live Data Simulation & Model Integration	ThingSpeak API integration, volume conversion (cylindrical), signal processing pipeline, live ML predictions, milestone3_predictions.csv	COMPLETE
4	Analytics Output & Dashboard Visualisation	3 structured CSVs (35 KPI columns), Power BI Desktop dashboard (8 visuals), DAX measures, Python auto-refresh, live_tanks.csv	COMPLETE
5	Result Analysis & Industry Interpretation	Final industry report, safety improvement quantification, inventory management ROI analysis, operational recommendations	UPCOMING

7.1 Complete File Artifact Registry

File	Format	Description
`smart_tank_dataset.csv`	CSV	Full 10,000-row synthetic dataset (22 columns)
`model_rf_regressor.pkl`	Joblib PKL	Trained Random Forest Regressor — fill level
`model_rf_classifier.pkl`	Joblib PKL	Trained Random Forest Classifier — risk label
`scaler.pkl`	Joblib PKL	Fitted StandardScaler
`label_encoder_risk.pkl`	Joblib PKL	Fitted LabelEncoder for risk class
`milestone3_predictions.csv`	CSV	Live prediction output (15 columns)
`milestone4_dashboard_data.csv`	CSV	Master dashboard source (35 KPI columns)
`milestone4_live_latest.csv`	CSV	Latest reading per tank for gauge/card visuals
`milestone4_daily_summary.csv`	CSV	Daily aggregates per tank
`live_tanks.csv`	CSV	Real-time feed refreshed every 15 seconds
`fetch_thingspeak.py`	Python	ThingSpeak data fetcher + volume converter
`auto_refresh_powerbi.py`	Python	Power BI Desktop auto-refresh via pyautogui

`SmartTank_M4_Dashboard.pbix`	Power BI	Complete 8-visual dashboard with DAX measures
`eda_dashboard.png`	PNG	6-panel EDA visualisation dashboard
`feature_importance.png`	PNG	Random Forest feature importance chart
`milestone3_dashboard.png`	PNG	6-panel live prediction dashboard

8. Conclusion

The Smart Storage Tank Level & Usage Prediction System has successfully delivered a production-grade IoT and machine learning pipeline across four completed milestones. The system addresses the critical operational challenges of petrochemical tank monitoring — including overflow risk, dry-run conditions, and inventory imbalance — through a five-layer architecture that combines physical sensing, embedded processing, cloud storage, ML inference, and live dashboard visualisation.

The Random Forest models trained on 10,000 synthetic records achieve exceptional performance: R^2 = 1.000000 for fill-level regression and 100% accuracy for three-class risk classification. In live deployment via the ThingSpeak pipeline, the system delivers predictions with an average error of 4.34 litres against tank capacities of over 50,000 litres — an accuracy rate exceeding 99.99%.

The Power BI Desktop dashboard provides plant operators with an immediate, actionable view of tank fill levels, risk alerts, consumption trends, and prediction accuracy through eight interactive visuals refreshed every 15 seconds via the Python auto-refresh solution. All operational data is traceable from the raw HC-SR04 ultrasonic distance reading through signal processing, volume conversion, ML inference, and dashboard display.

Milestone 5 will consolidate the results of all four completed milestones into a final industry interpretation report, quantifying the safety improvements achieved, the inventory management benefits enabled, and the economic value delivered to L&T Construction plant operations.

Project Status: Milestones 1–4 Complete

Full end-to-end pipeline operational — ESP32 sensor → ThingSpeak → ML inference → Power BI live dashboard